

# IPv4

vs.

# IPv6

Deployed 1981

32-bit IP address

4.3 billion addresses

Addresses must be reused and masked

Numeric dot-decimal notation

**192.168.5.18**

DHCP or manual configuration

Deployed 1998

128-bit IP address

$7.9 \times 10^{28}$  addresses

Every device can have a unique address

Alphanumeric hexadecimal notation

**50b2:6400:0000:0000:6c3a:b17d:0000:10a9**

(Simplified - 50b2:6400::6c3a:b17d:0:10a9)

Supports autoconfiguration

# IPV6 Address

128-Bits

2001:4860:4860:0000:0000:0000:0000:8844

48-Bits

16-Bits

64-Bits

Network part

Subnet ID

Client ID



2001:0000:3238:DFE1:63:0000:0000:FEFB

### IPv4-compatible IPv6 address




### IPv4-mapped IPv6 address

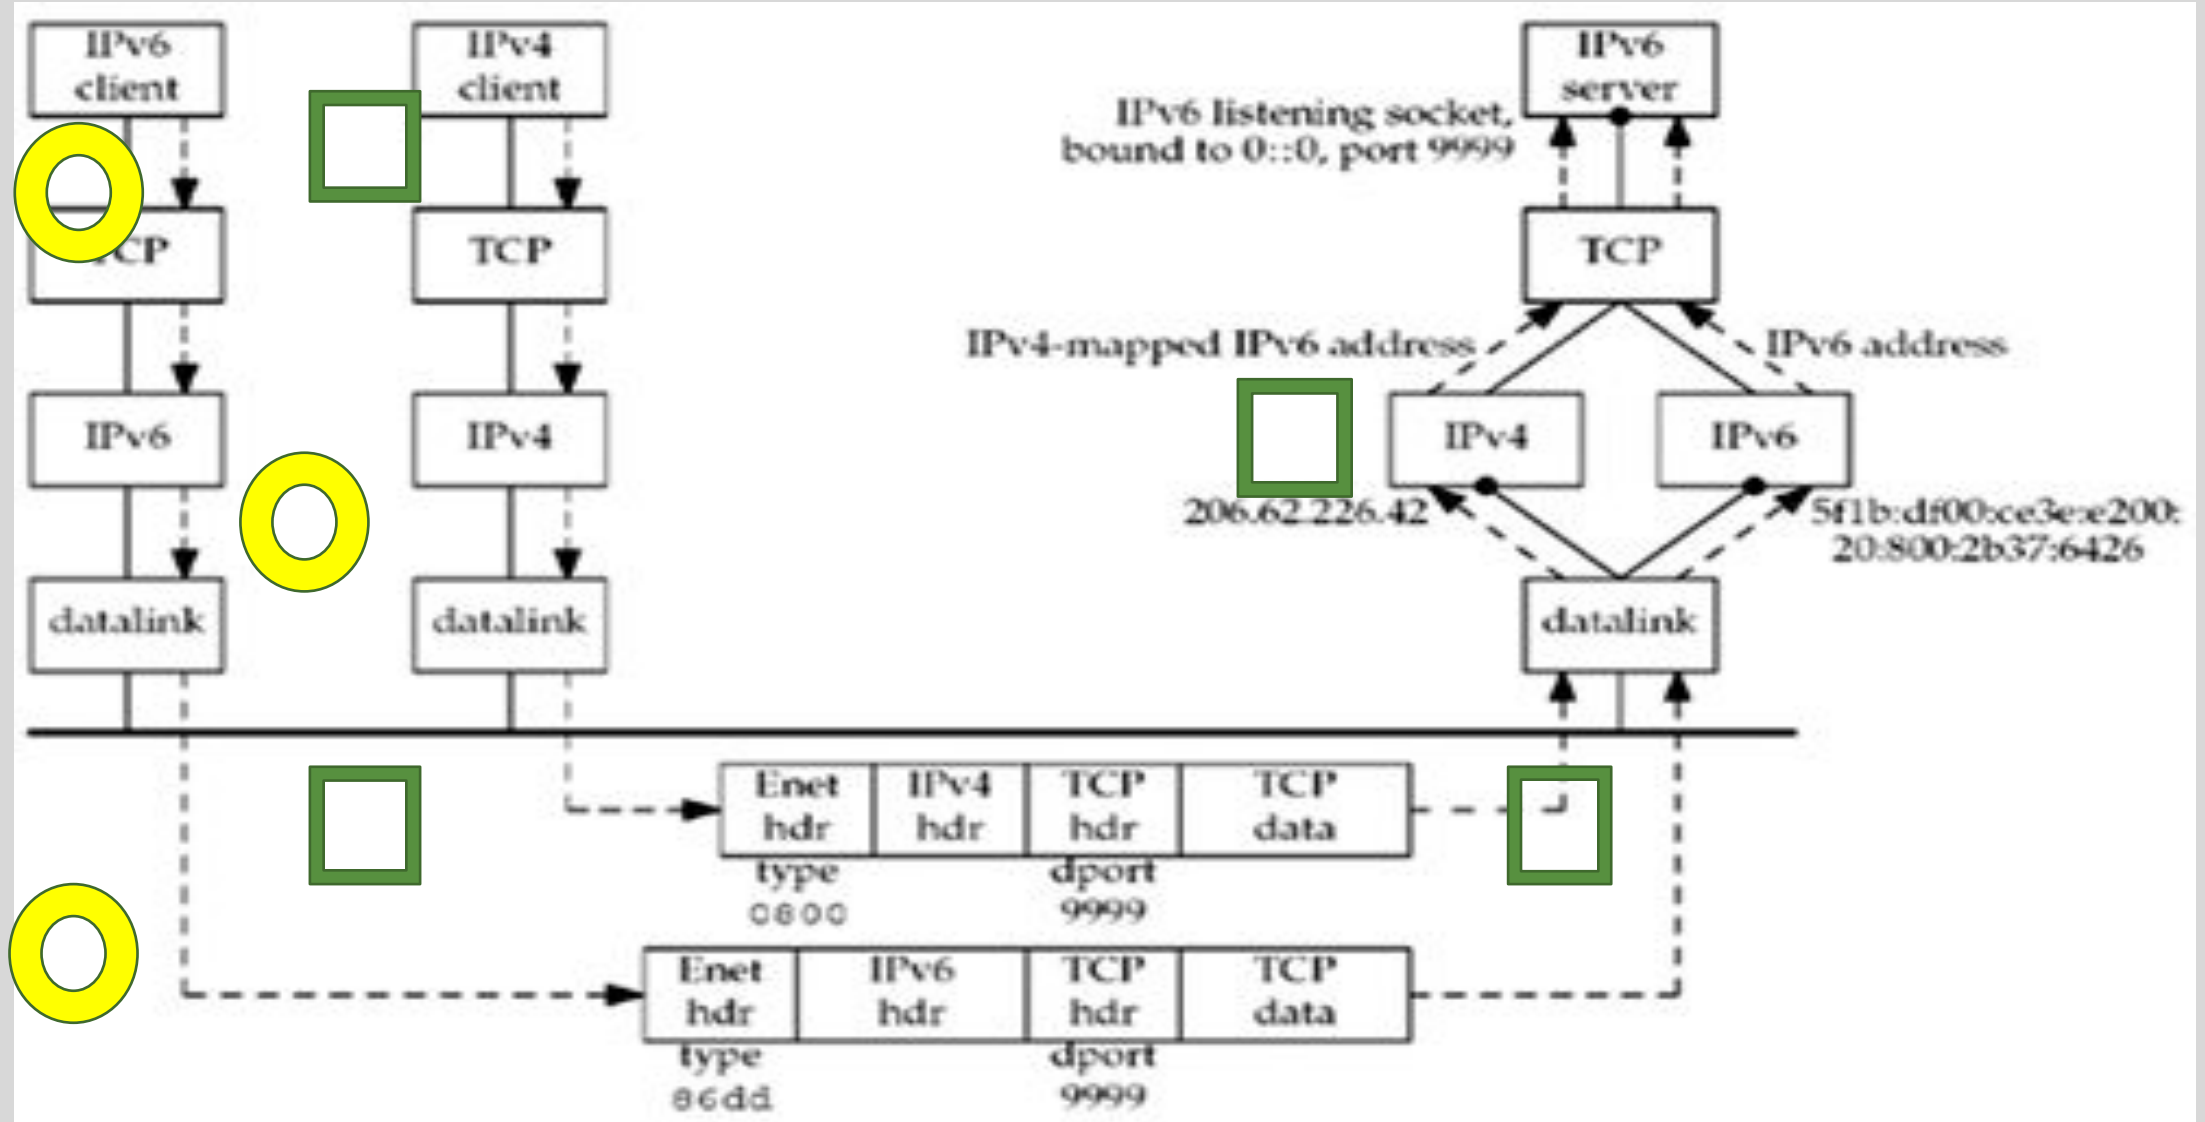


# IPV4 and IPV6 Interoperability

- Transition will be gradual from IPV4 to IPV6
- During this transition phase, existing IPV4 applications must continue to work with newer IPV6 Applications



# IPv4 Client and IPv6 Server



# IPV4 Client and IPV6 Server

**Step 1. IPV6 Server starts, creates a listening IPV6 socket and it binds wildcard address to the socket**

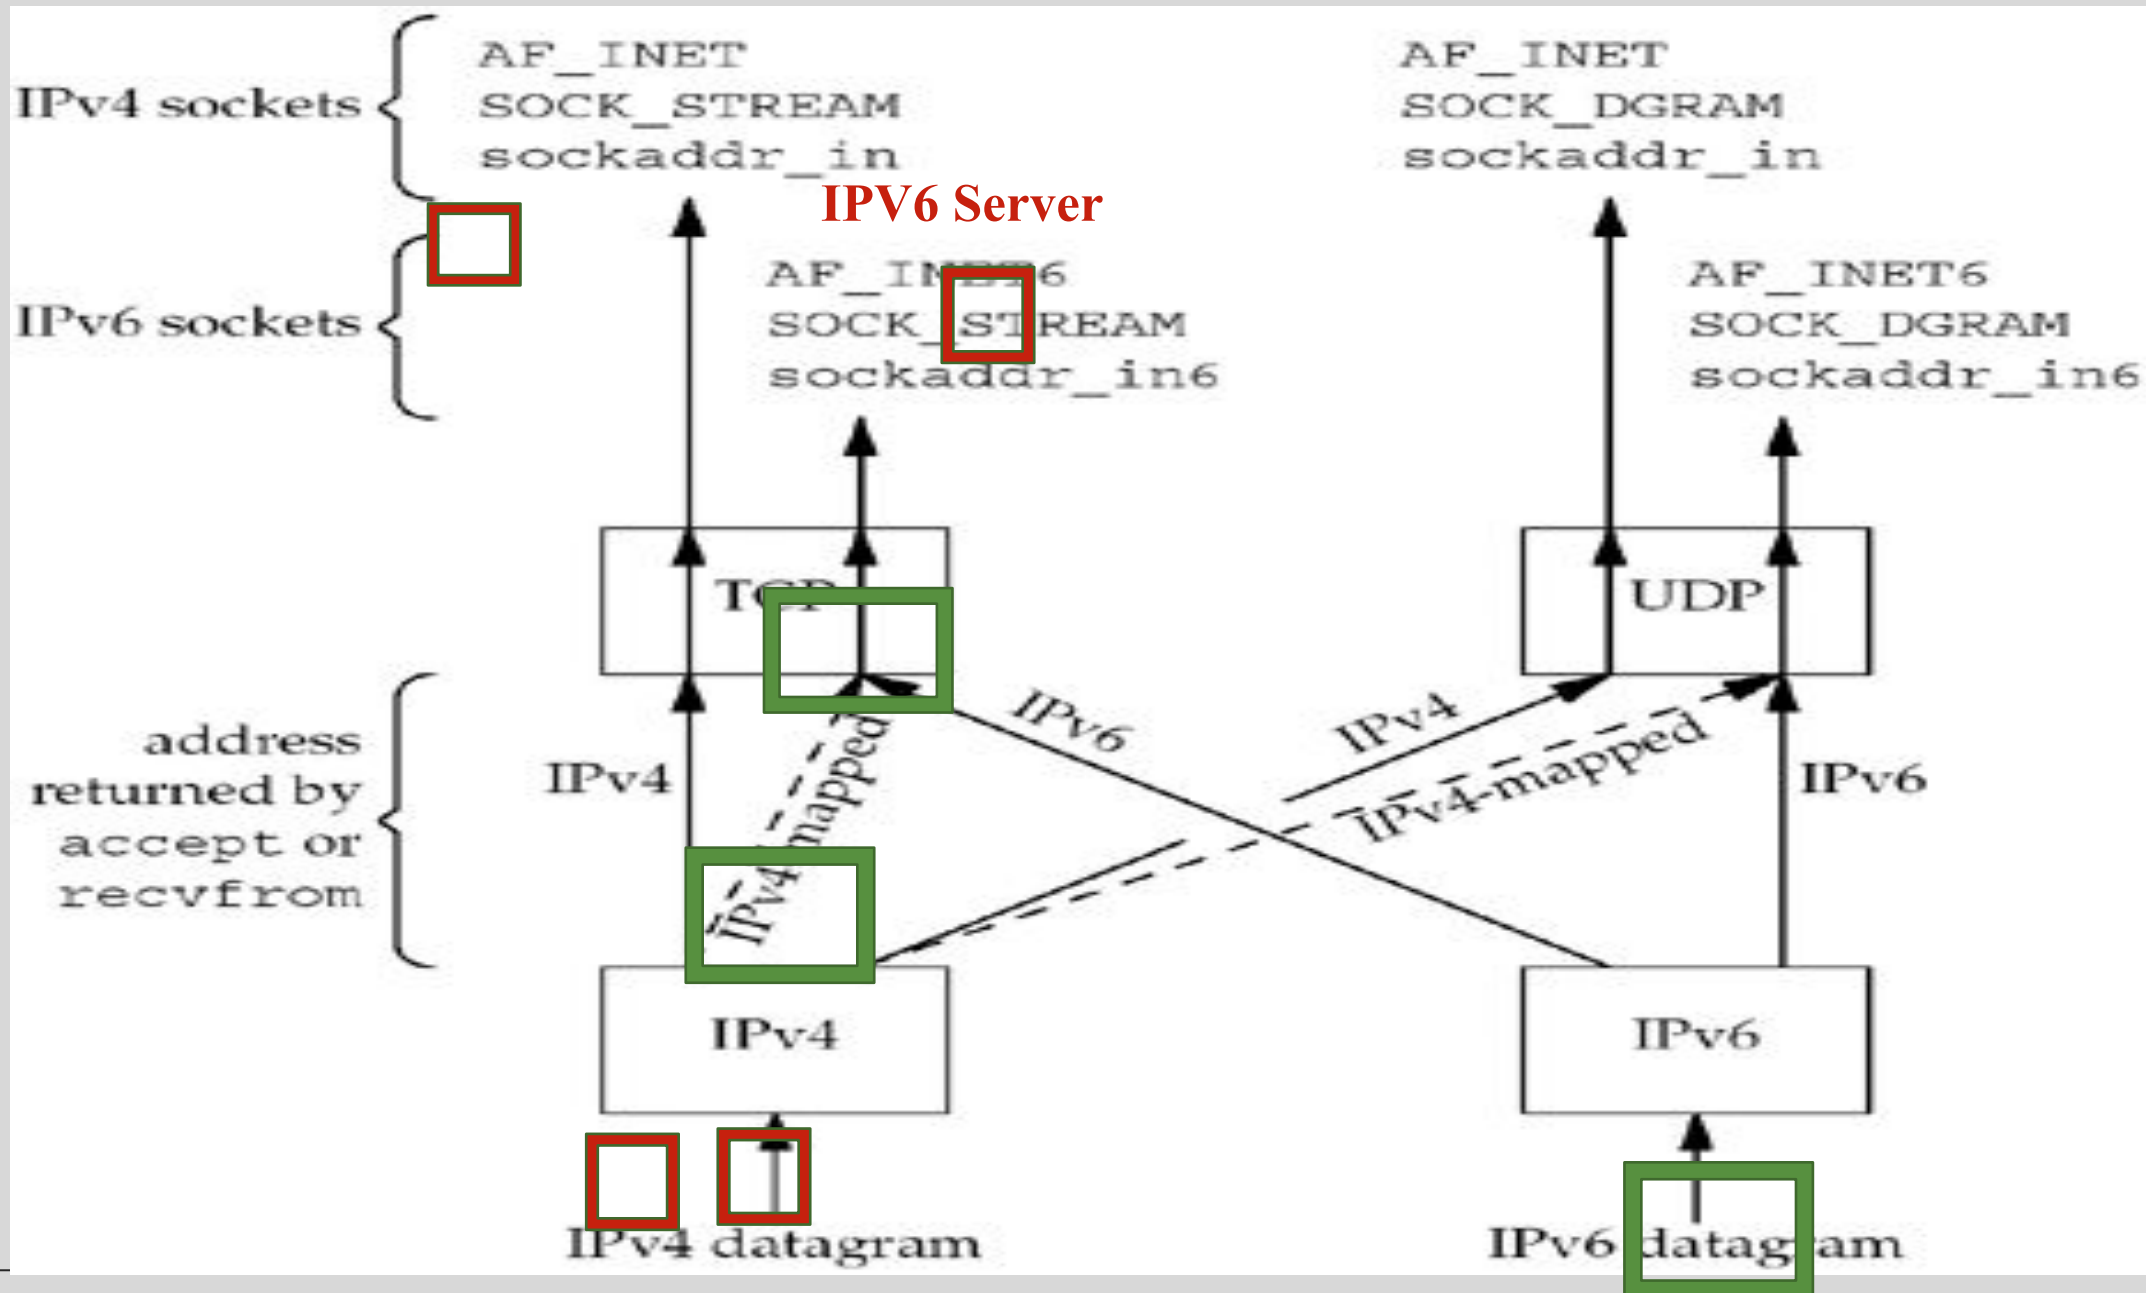
**Step 2. IPV4 client calls gethostbyname and finds an A record..**

**Steps 3. The Client calls connect and client's host sends an IPV4 SYN to server**

**Step 4. The Server host receives IPV4 SYN directed to IPV6 socket, sets a flag indicating the this connection is using IPV4 mapped IPV6 address and responds with IPV4 SYN/ACK**

**Step 5. When the server host sends IPV4-mapped-IPV6 address, IP Stack generates IPV4 datagram to IPV4 address**

**Step 6. Dual Stack handles all the finer details and the Server is unaware that it is communicating with IPV4 client**



# IPv6 Client – IPv4 Server

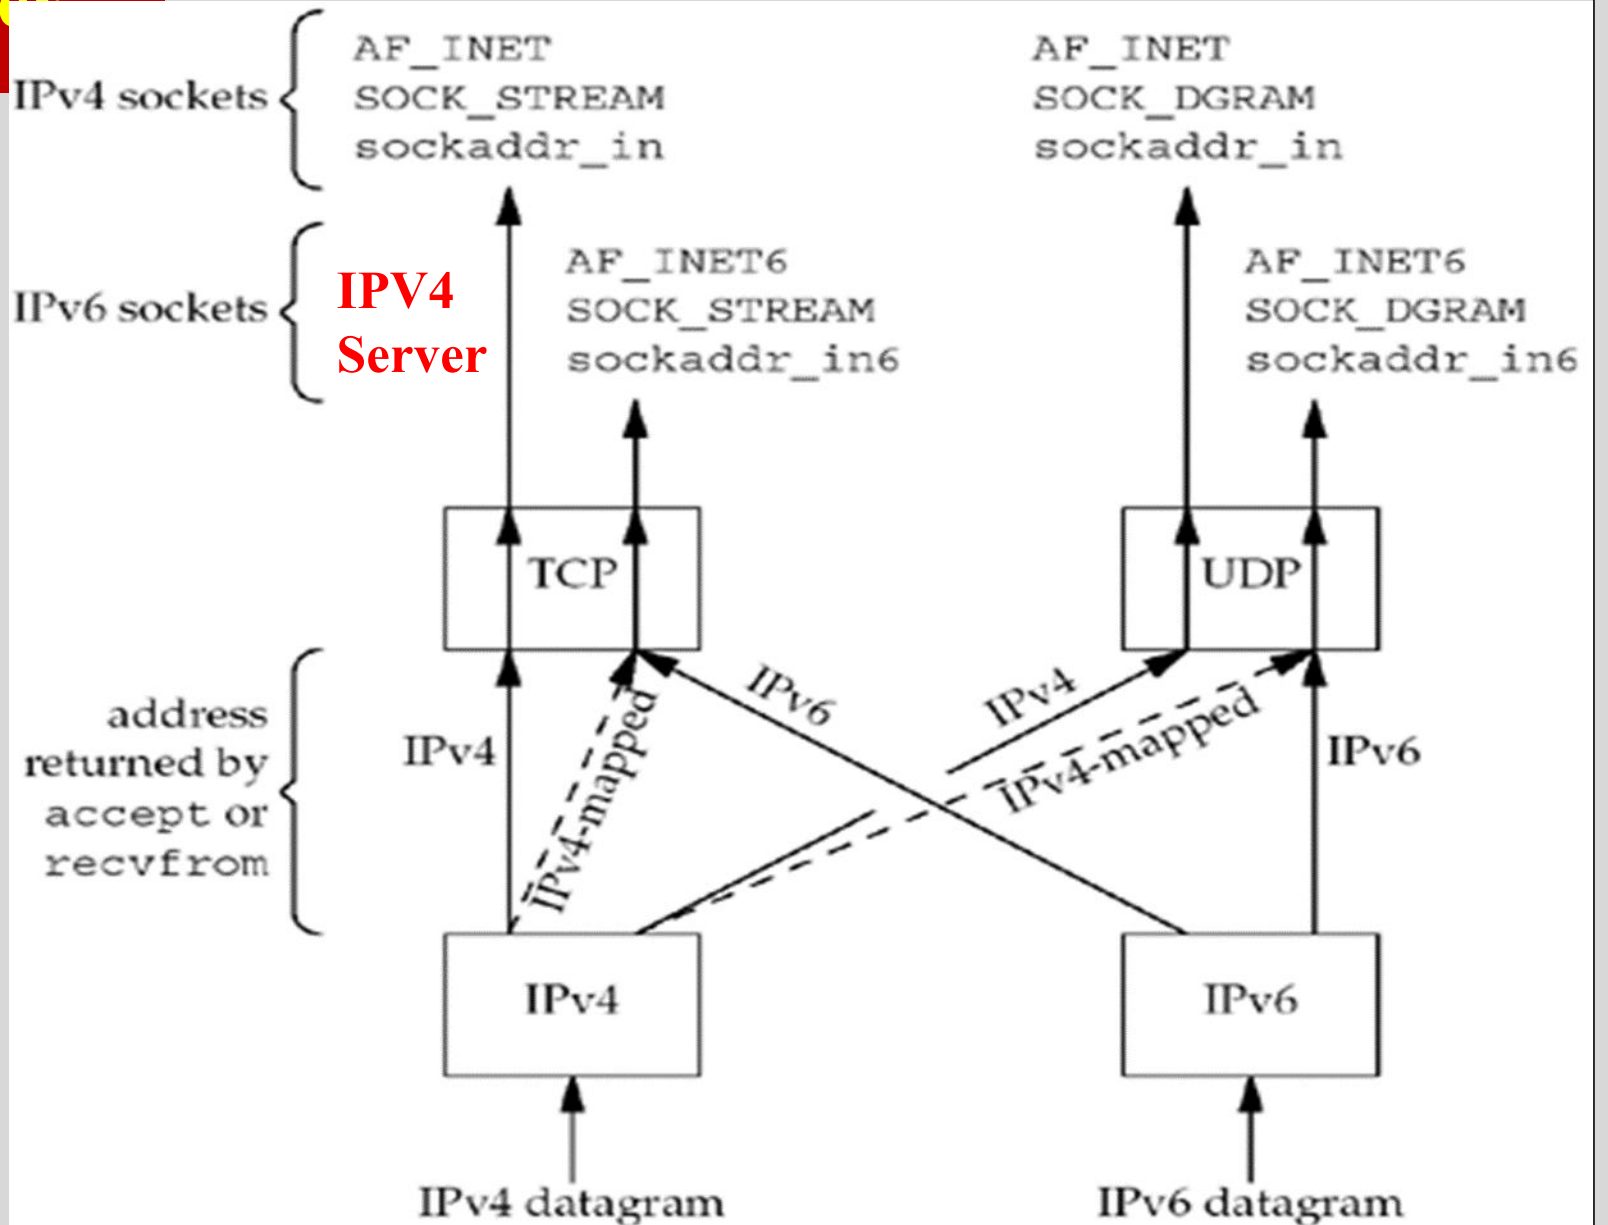
- An IPv4 server starts on an IPv4-only host and creates an IPv4 listening socket.
- The IPv6 client starts and calls `getaddrinfo` asking for only IPv6 addresses
- it requests the `AF_INET6` address family and sets the `AI_V4MAPPED` flag in its hints structure.
- Since the IPv4-only server host has only A records, we see that an IPv4-mapped IPv6 address is returned to the client.
- The IPv6 client calls `connect` with the IPv4-mapped IPv6 address in the IPv6 socket address structure.



# IPv6 Client – IPv4 Server

- **The kernel detects the mapped address and automatically sends an IPv4 SYN to the server.**
- **The server responds with an IPv4 SYN/ACK, and the connection is established using IPv4 datagrams.**

# IPV6 Client – IPV4 Server



# IPv4 Client - IPv6 Server

- When the server host sends to the IPv4-mapped IPv6 address, its IP stack generates IPv4 datagrams to the IPv4 address. Therefore, all communication between this client and server takes place using IPv4 datagrams.
- Unless the server explicitly checks whether this IPv6 address is an IPv4-mapped IPv6 address (using the `IN6_IS_ADDR_V4MAPPED` macro described in Section 12.4), the server never knows that it is communicating with an IPv4 client.
- The dual-protocol stack handles this detail. Similarly, the IPv4 client has no idea that it is communicating with an IPv6 server.

## IPv6 Address Macro, Function, Option

- IPv6 address testing macros:
  - IN6\_IS\_ADDR\_\* (e.g. V4MAPPED)
- Protocol independent socket address functions:
  - sock\_\* (e.g. cmp\_addr) (section 3.8)
- IPv6\_ADDRFORM socket option:
  - change a socket type between IPv4 and IPv6, by setsockopt function with IPv6\_ADDRFORM option

# IPV6 Address Testing Macros

- There are a small class of IPV6 Applications, that must know whether they are communicating with IPV4 peers.
- These applications need to know whether the peer is using IPV4 mapped IPV6 address
- There are twelve macros that help in this requirement

# IPV6 Address Testing Macros

```
int IN6_IS_ADDR_V4MAPPED(const struct in6_addr *aptr);
```

```
int IN6_IS_ADDR_V4COMPAT(const struct in6_addr *aptr);
```

# Source Code Portability

**IPV4 Applications can be converted to IPV6 with minor modifications**

**But we need to check if the receipient host has support for IPV6**

**We may do this using #ifdefs in the code.. But code will have too many of these #ifdefs**

**Better way is to make the program protocol independent**

# Source Code Portability

The first step is... remove all `gethostbyname` and `gethostbyaddr` and use

`getaddrinfo` and `getnameinfo` functions


This helps deal with `sockaddr` as opaque objects.

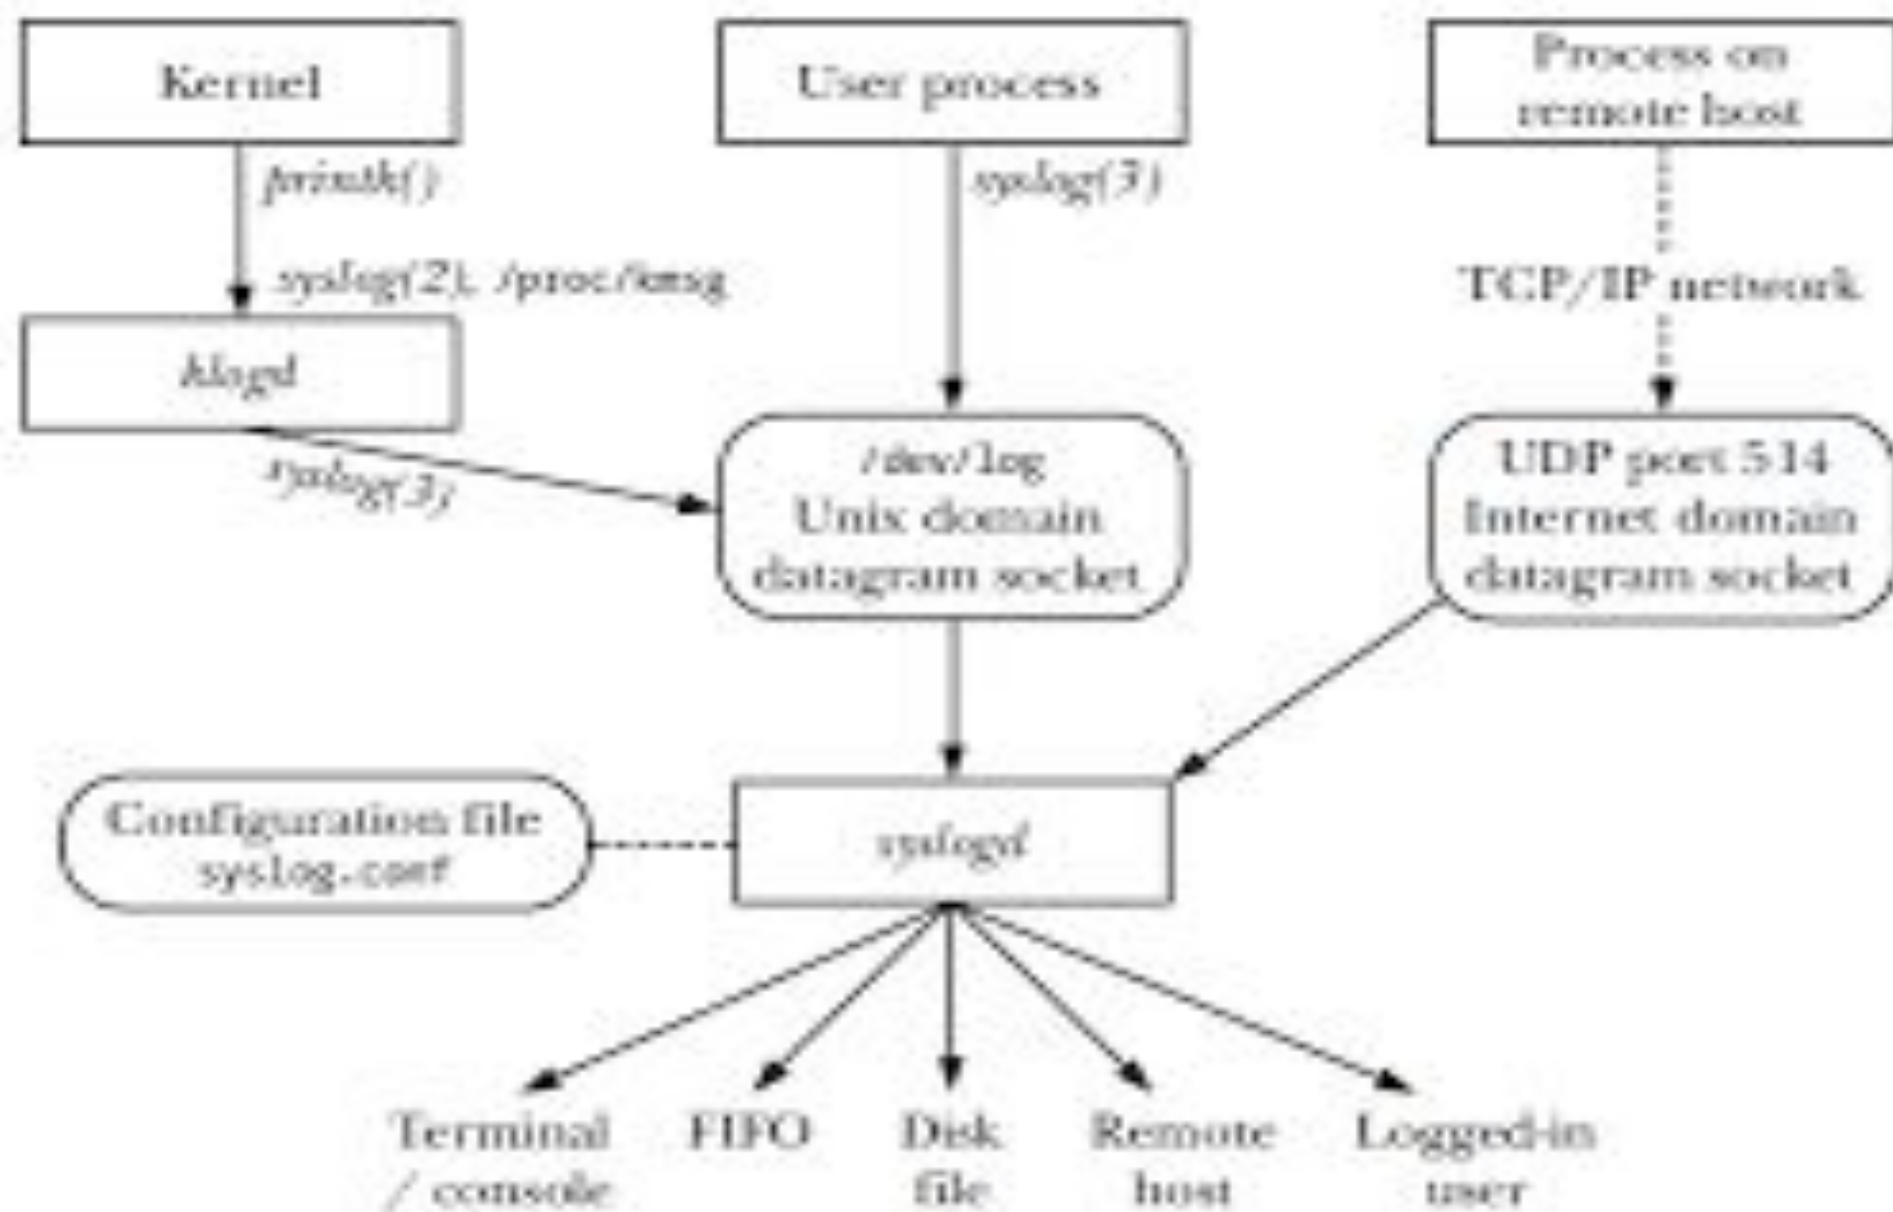


# Syslogd Deamon

## Syslogd Daemon

- Different ways to communicate with syslogd
  - Unix domain socket
  - UDP socket (on port 514)
  - A file opened for accepting messages from kernel





Syslog stands for System Logging Protocol and is a standard protocol **used to send system log or event messages to a specific server**, called a syslogserver.

It is primarily used to collect various device logs from several different machines in a central location for monitoring and review.

**void syslog(int priorityLevel, char \*msg);**

## 12.3 syslog function

- Log message have a level between 0 and 7.

<i>level</i>	value	description
LOG_EMERG	0	system is unusable ( highest priority )
LOG_ALERT	1	action must be taken immediately
LOG_CRIT	2	critical conditions
LOG_ERR	3	error conditions
LOG_WARNING	4	warning conditions
LOG_NOTICE	5	normal but significant condition (default)
LOG_INFO	6	informational
LOG_DEBUG	7	debug-level message ( lowest priority )

Figure 12.1 *level* of log message.

## Use of syslog() function

```
static void
err_doit (int errnoflag, int level, const char *fmt, va_list ap)
{
    int          errno_save, n;
    char  buf[MAXLINE];

    errno_save = errno;
    vsprintf(buf, fmt, ap); n = strlen(buf);
    if (errnoflag)
        snprintf(buf+n, sizeof(buf)-n, ": %s", strerror(errno_save));
    strcat(buf, "\n");

    if (daemon_proc) {
        syslog (level, buf); /* send to syslogd daemon for logging */
    } else {
        fflush(stdout);      /* print out to the screen */
        fputs(buf, stdout);
        fflush(stdout);
    }
    return;
}
```



# Daemon Init

```
#define MAXFD 64
extern int daemon_proc; /* defined in error.c */
void daemon_init (const char *pname, int facility)
{
    int i;
    pid_t pid;
    if ( (pid = Fork()) != 0)
        exit(0); /* parent terminates */
    /* 1st child continues */
    setsid(); /* become session leader */
    Signal(SIGHUP, SIG_IGN);
    if ( (pid = Fork()) != 0)
        exit(0); /* 1st child terminates */
    /* 2nd child continues */
    daemon_proc = 1; /* for our err_XXX() functions */
    chdir("/"); /* change working directory */
    umask(0); /* clear our file mode creation mask */
    for (i = 0; i < MAXFD; i++)
        close(i);
    openlog(pname, LOG_PID, facility);
}
```

## Daytime tcp server as a daemon

```
int main(int argc, char **argv)
{
    int listenfd, connfd; socklen_t addrlen, len;
    struct sockaddr_in *cliaddr; char buff[MAXLINE]; time_t ticks;

    daemon_init (argv[0], 0);
    if (argc == 2)
        listenfd = Tcp_listen(NULL, argv[1], &addrlen);
    else if (argc == 3)
        listenfd = Tcp_listen(argv[1], argv[2], &addrlen);
    else
        err_quit("usage: daytimetcpsrv2 [ <host> ] <service or port>");
    cliaddr = Malloc(addrlen);
    for ( ;; ) {
        len = addrlen;
        connfd = Accept(listenfd, cliaddr, &len);
        err_msg("connection from %s", Sock_ntop(cliaddr, len));
        ticks = time(NULL);
        snprintf(buff, sizeof(buff), "%.24s\r\n", ctime(&ticks));
        Write(connfd, buff, strlen(buff));
        Close(connfd);
    }
}
```

# inetd daemon

## *internet superserver*

There are a lot of internet servers that will be running on a typical Unix system

- FTP, telnet, rlogin, TFTP, daytime, time, echo Servers

In earlier systems, each service had a process (daemon) associated with.

- each process does the *same startup tasks*  
creating sockets, binding listening....
- each process took *an entry in the process table*
- these processes are *asleep most of the time* waiting for requests.

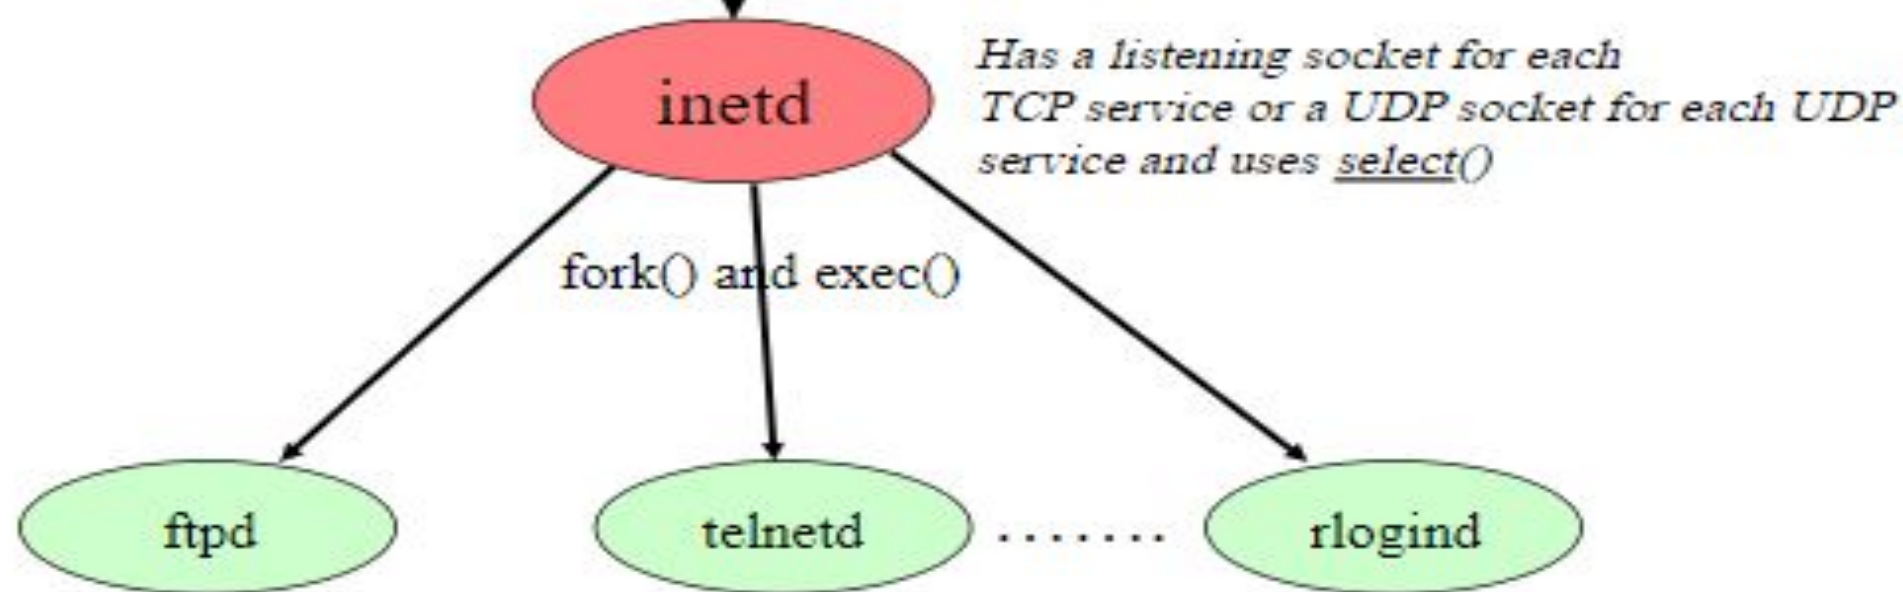
After 4.3 BSD Unix system, a single server **inetd** daemon is used to wait for requests for all of these services and then create the corresponding server process on demand.



### */etc/inetd.conf* file

ftp	stream	tcp	nowait	root	/usr/sbin/in.ftpd	in.ftpd
telnet	stream	tcp	nowait	root	/usr/sbin/in.telnetd	in.telnetd
name	dgram	udp	wait	root	/usr/sbin/in.tnamed	in.tnamed
shell	stream	tcp	nowait	root	/usr/sbin/in.rshd	in.rshd
login	stream	tcp	nowait	root	/usr/sbin/in.rlogind	in.rlogind
tftp	dgram	udp	wait	root	/usr/sbin/in.tftpd	in.tftpd -s /tftpboot

*reads the configuration file upon startup*



## steps performed by `inetd`

